

RAG HALLUCINATION POLICY

01 PRIMARY SOURCE RULE

When a student uploads study material, the AI system **must use the retrieved content** from that material as the primary source for answering all document-based questions. No external prior knowledge may supersede direct source evidence.

02 NO GUESSING RULE

The AI must **never guess missing information**. If the retrieved material does not contain enough context, it must strictly return the standardized response:

"I couldn't find enough information in the provided study material to answer this confidently."


03 STRICT NON-FABRICATION POLICY

The AI model is strictly prohibited from inventing or extrapolating any of the following parameters:

× Facts	× Statistics	× Definitions	× Formulas
× Dates	× Names	× Quotes	× References
× Page Numbers	× Citations	× URLs	× PDF Claims

04 RETRIEVAL REQUIREMENT ARCHITECTURE

All student queries must strictly flow through the sequential retrieval and evaluation multi-agent pipeline:



05 RELEVANCE THRESHOLD STANDARD

If retrieved chunks do not cross the minimum relevance confidence threshold, the system **must not force an answer** from the document. It must gracefully decline and inform the student that sufficient matching material was not found.

06 TUTOR AGENT DIRECTIVE POLICY

SYSTEM DIRECTIVE RULES FOR TUTOR AGENT

- You are an educational tutor. Use retrieved study material as primary evidence.
- Answer **only** what can be strictly supported by the retrieved context and question.
- Do not invent facts, citations, quotations, page numbers, references, or content.
- If the retrieved context is insufficient, explicitly state the fallback response.
- Simplification for academic level is permitted, provided source meaning remains unaltered.
- Never claim unsupported information originated from the uploaded study material.

07 RESEARCH AGENT QUALITY CONTROL

MANDATORY VERIFICATION AUDIT

- Cross-check Tutor Agent's drafted response against source context.
- Flag and eliminate any unsupported claims or hallucinatory extensions.
- Identify critical missing context and correct factual discrepancies.
- Improve clarity and educational structure.
- **Strict Rule:** Never inject unsupported details merely to elongate the answer.

08 EVALUATOR AGENT FINAL AUDIT CHECKLIST

✓ Answer supported by context?	✓ Zero unsupported claims?
✓ Facts fully verified?	✓ No fake citations created?
✓ Source meaning preserved?	✓ Academic level appropriate?
✓ Flagged claims removed?	✓ Missing context acknowledged?

09 FINAL ANSWER RULE & IMPLEMENTATION PRINCIPLE

System Preference Priority: The AI architecture must systematically prefer returning:

"I don't have enough information from the provided material."

over speculative phrasing such as:

"I think the answer is..."

IMPLEMENTATION NOTE

A RAG policy alone cannot guarantee zero hallucinations. Operational success requires combining **Retrieval + Relevance Checking + Grounded Prompts + Multi-Agent Evaluation + Source Traceability + Systematic Testing.**